

A Reinforcement Learning Framework for Energy-Efficient and Grid-Resilient Residential HVAC Control in Tropical Climates

Mahjabeen Rahman

2023-1-60-253

*Department of Computer Science and Engineering
East West University, Dhaka, Bangladesh*

Sheikh Tanvir Mohiuddin

2022-1-60-351

*Department of Computer Science and Engineering
East West University, Dhaka, Bangladesh*

Supervised by: Ahmed Abdal Shafi Rasel

*Department of Computer Science and Engineering
East West University, Dhaka, Bangladesh
ahmed.shafi@ewubd.edu*

Monika Asha Momo

2023-1-60-180

*Department of Computer Science and Engineering
East West University, Dhaka, Bangladesh*

Habibur Rahman

2023-1-60-024

*Department of Computer Science and Engineering
East West University, Dhaka, Bangladesh*

Abstract—Residential HVAC systems account for approximately 40% of household electricity consumption globally, yet most intelligent control frameworks assume uninterrupted power supply and temperate climates—two assumptions that do not hold in Bangladesh, where unscheduled load-shedding is routine and outdoor temperatures routinely exceed 35°C. This paper presents and evaluates a Soft Actor-Critic (SAC) based reinforcement learning framework for energy-efficient, thermally comfortable, and grid-resilient residential HVAC control in a tropical climate. The reward function jointly penalises energy consumption, deviations from Fanger’s Predicted Mean Vote (PMV) comfort index, rapid control switching that causes compressor wear, and grid-instability risk. We train and evaluate three single-agent algorithms — SAC, Proximal Policy Optimisation (PPO), and Twin Delayed DDPG (TD3) — and one multi-agent algorithm (MADDPG) for five-zone coordination, using the Sinergym/EnergyPlus simulation platform with a Dhaka-specific `dhaka-radiant-residential-v1` environment. Over 5000 evaluation steps, SAC achieves the lowest HVAC energy consumption (487.3 kWh), the smallest comfort-violation rate (22.6%), and the lowest switching frequency (0.0623), yielding a multi-objective score of 0.847 against 0.612 for TD3 and 0.389 for PPO. A benchmark comparison under identical reward structures demonstrates that SAC’s entropy-regularised exploration, rather than the actor-critic architecture itself, is the primary driver of superior performance. MADDPG achieves the lowest peak grid-risk signal (0.68) through decentralised zone coordination, but incurs higher switching frequency across agents. Wilcoxon signed-rank tests ($p < 0.001$) confirm that all pairwise differences between SAC and the other controllers are statistically significant.

Index Terms—Soft Actor-Critic, HVAC control, thermal comfort, PMV, load-shedding resilience, Sinergym, reinforcement learning, green computing, multi-agent RL, tropical climate, Bangladesh

1. Introduction

1.1. Background and Motivation

Heating, Ventilation, and Air Conditioning (HVAC) systems are the dominant energy load in residential buildings worldwide, consuming roughly 40% of total household electricity [1]. As Bangladesh’s ICT sector grows at 15–20% per year and urban heat islands intensify, residential cooling demand is forecast to double by 2030 [2]. Deep Reinforcement Learning (DRL) has emerged as a compelling alternative to rule-based HVAC control because DRL agents learn optimal policies directly from environment interaction without requiring explicit thermodynamic models [3], [4].

Despite rapid progress, three limitations undermine the practical deployability of existing frameworks. First, nearly all studies focus on commercial offices with predictable occupancy schedules [5], [4]; residential homes exhibit far higher randomness [6]. Second, the overwhelming majority of published HVAC-RL work assumes a continuously available electricity supply — an assumption that fails across large parts of South Asia and sub-Saharan Africa. Third, most frameworks treat the AI scheduler’s own computational footprint as negligible, ignoring the fact that heavy neural-network inference can itself consume meaningful energy [7].

1.2. Problem Statement

Standard rule-based thermostats suffer from four interrelated deficiencies [4], [6]: (i) binary ON/OFF logic causes temperature overshoot and energy waste; (ii) rapid control switching inflicts mechanical wear on compressors; (iii) temperature-only control ignores humidity and activity factors that determine true comfort; and (iv) existing RL controllers are trained under the assumption of a stable grid, leaving them incapable of proactive pre-cooling before a load-shedding event.

1.3. Contributions

This paper makes four concrete contributions:

- 1) A **PMV-and-grid-aware SAC** controller trained on the `dhaka-radiant-residential-v1` environment, the first publicly available Dhaka-specific Sinergym benchmark.
- 2) A **benchmark comparison** of SAC, PPO, and TD3 under an identical reward structure, isolating the contribution of entropy-regularised exploration from the actor-critic architecture.
- 3) A **multi-agent MADDPG** extension for five-zone coordination that reduces peak grid-risk signal at the cost of higher per-agent switching frequency.
- 4) A **robustness analysis** showing how comfort-violation rates degrade as measurement noise increases, revealing SAC’s superior noise tolerance.

2. Literature Review

2.1. Foundational RL Approaches for HVAC Control

Early work established continuous action spaces as superior to discrete ones for fine-grained setpoint tracking [4]. Chen *et al.* [8] combined differentiable Model Predictive Control (MPC) with RL to improve sample efficiency (Gnu-RL), while Ding *et al.* [9] introduced MB2C, a model-based deep RL controller for multi-zone buildings that predicts future dynamics at the cost of higher computational complexity. Mnih *et al.* [10] established DQN as a foundational benchmark; Lillicrap *et al.* [11] introduced DDPG for continuous control; and Haarnoja *et al.* [12] proposed SAC with a maximum-entropy objective that proves critical in this study.

2.2. Residential and Comfort-Focused Approaches

Gao *et al.* [13] jointly optimised energy and thermal comfort with DeepComfort. Ding *et al.* [14] addressed multi-zone residential HVAC, and Boutahri & Tilioua [15] validated an RL controller using BOPTEST, achieving an 8.8% energy reduction over a rule-based baseline in Morocco. However, all of these studies remain tied to a single, stable-grid climate context. Salimi *et al.* [7] highlighted the underexplored role of building occupancy uncertainty in RL training, a challenge directly relevant to Dhaka’s variable household patterns.

2.3. SAC-Based and Context-Adaptive Methods

Coraci *et al.* [16] first demonstrated SAC’s viability for online building control. Wang *et al.* [17] showed that outdoor-temperature-dependent reward weighting improves SAC’s performance across climates. Alshammari *et al.* [18] introduced SACEM, a SAC-plus-Cross-Entropy-Method hybrid prioritising comfort under time-of-use pricing. Lim *et al.* [19] improved personalised comfort robustness under building-parameter uncertainty. Hu *et al.* [20] proposed GASAC — a generative-adversarial variant that learns from expert demonstrations without an explicit reward function. Yu *et al.* [21] demonstrated that tropical cities behave categorically differently from temperate ones under RL control, directly motivating a Dhaka-specific evaluation.

2.4. Multi-Agent and Grid-Aware Approaches

Vázquez-Canteli & Nagy [22] reviewed RL for demand response, identifying grid-interaction as an open challenge. Wang & Hong [23] surveyed building controls with RL and flagged load-shedding resilience as an understudied problem. Table 1 summarises the gap each prior study leaves unaddressed.

TABLE 1. PRIOR SAC-BASED HVAC STUDIES AND UNADDRESSED GAPS

Study	Method	Gap Left Unaddressed
Coraci <i>et al.</i> [16]	SAC	Not residential; no switching penalty
Ding <i>et al.</i> [14]	RL (multi-zone)	No switching penalty; stable-grid assumption
Boutahri & Tilioua [15]	RL+BOPTEST	Mediterranean climate only; no grid outage
Wang <i>et al.</i> [17]	Adaptive SAC	No occupancy or grid-context weighting
Alshammari <i>et al.</i> [18]	SAC+CEM	Not residential; no switching penalty
Lim <i>et al.</i> [19]	SAC personal.	Not residential; no switching penalty
Hu <i>et al.</i> [20]	Gen-Adv SAC	Needs expert demos; no grid-outage context
Yu <i>et al.</i> [21]	RL+urban clim.	City-level; not household-scale

3. Research Questions and Hypotheses

RQ1: Can a continuous SAC controller outperform standard ON/OFF thermostats in reducing residential HVAC energy during both summer and winter in a tropical climate?

RQ2: Does a switching penalty in the reward function reduce equipment wear without sacrificing indoor comfort?

RQ3: Does a PMV-based comfort model outperform temperature-only control under random daily occupancy disturbances?

RQ4: Does incorporating a grid-reliability context signal improve realised comfort during simulated load-shedding events versus a grid-unaware baseline?

RQ5 (Benchmark): Does entropy-regularised exploration (SAC’s maximum-entropy objective) drive the performance advantage, or would any actor-critic method achieve comparable results under an identical reward structure?

Corresponding hypotheses: H1 — SAC reduces HVAC energy by $\geq 10\%$ over rule-based control; H2 — the switching penalty reduces Δf_{switch} by $\geq 30\%$ without increasing PMV violation; H3 — PMV-based control reduces comfort violations by $\geq 15\%$ versus temperature-only; H4 — grid-aware SAC maintains PMV within the comfort band for $\geq 20\%$ more timesteps during outages than grid-unaware SAC; H5 — SAC significantly outperforms PPO and TD3 under identical rewards, with the gap attributable to entropy-driven exploration ($p < 0.001$, Wilcoxon).

4. Methodology

4.1. Simulation Environment

All experiments use **Sinergym 3.4.0** [24], a Python reinforcement-learning testbed built atop EnergyPlus 23.1.0. The environment `dhaka-radiant-residential-v1` models a single-family residential house with five thermal zones (Living, Kitchen, Bedroom01–03) in Dhaka, Bangladesh, driven by a Dhaka .epw weather file (tropical monsoon, 25–38°C summer peak). Radiant heating and cooling panels are controlled independently in the MADDPG experiments; SAC, PPO, and TD3 receive the full multi-zone state and output setpoints aggregated across zones.

4.2. State Space

The 29-dimensional observation vector includes: zone air temperatures (T_z), relative humidity (ϕ_z), outdoor dry-bulb temperature (T_{out}), solar irradiance, total facility electricity demand (kW), time-of-day encoding, occupancy indicator (ρ), and the grid-reliability signal $\gamma_{\text{grid}} \in [0, 1]$ — the estimated probability of an imminent load-shedding event derived from a historically calibrated hourly outage schedule for Dhaka’s distribution network.

4.3. Action Space

Single-agent controllers (SAC, PPO, TD3) output a single normalised setpoint $a_t \in [-1, +1]$, linearly mapped to the cooling/heating setpoint range [18, 28]°C. MADDPG operates five independent agents — one per zone — each producing $a_{t,z} \in [-1, +1]$.

4.4. Reward Function

At each timestep t the reward for single-agent controllers is:

$$R_t = -[w_1 P_{\text{energy}}(t) + w_2 |\text{PMV}(t)| + w_3 \Delta f_{\text{switch}}(t) + w_4 P_{\text{grid}}(t)] \cdot \Phi(s_{\text{context}}) \quad (1)$$

where $P_{\text{energy}}(t)$ is instantaneous HVAC power (kW), $|\text{PMV}(t)|$ is the absolute PMV deviation from neutral,

$\Delta f_{\text{switch}}(t) = |a_t - a_{t-1}|$ penalises rapid changes, and $P_{\text{grid}}(t) = \gamma_{\text{grid}}(t) \cdot |a_t|$ penalises high-power actions when grid instability is forecast. Weights are $w_1 = 0.30$, $w_2 = 0.35$, $w_3 = 0.20$, $w_4 = 0.15$. $\Phi(s_{\text{context}})$ zeroes the comfort penalty when $\rho = 0$ (unoccupied) and doubles w_4 when $\gamma_{\text{grid}} > 0.7$.

PMV proxy: following the simplified ISO 7730 Fanger formulation [25]:

$$\text{PMV}(t) \approx 0.25 (T_z - 24) + 0.01 (\phi_z - 50) \quad (2)$$

with velocity $v = 0.1$ m/s, metabolic rate $\text{met} = 1.1$, clothing $\text{clo} = 0.5$. A full ISO 7730 implementation is also included in the notebook for per-step PPD calculation.

4.5. Algorithm Configurations

SAC [12]: off-policy, maximum-entropy actor-critic with automatic temperature tuning; learning rate 10^{-3} ; replay buffer 10^5 ; batch size 256; discount $\gamma = 0.99$; trained for 1 000 episodes.

PPO [26]: on-policy proximal clipping ($\epsilon = 0.2$); 64 steps per rollout; entropy coefficient 0.01; learning rate 3×10^{-4} .

TD3 [27]: deterministic actor-critic with delayed target updates (policy delay $d = 2$); action noise $\sigma = 0.1$; clipping $c = 0.5$; learning rate 10^{-3} .

MADDPG [28] (5 agents): centralised critic, decentralised actors; per-agent replay buffer; batch size 32; 50 gradient updates post warm-up.

All single-agent models are trained via Stable-Baselines3; MADDPG via AgileRL. Hyperparameters are locked at deterministic defaults (random seed 42) to ensure reproducibility. Models are loaded and evaluated at 5 000 environment steps.

4.6. Evaluation Protocol

Each algorithm is evaluated for exactly 5 000 steps on the held-out test partition of the simulation (no retraining). Metrics recorded per step: HVAC power draw (kW), zone temperature, relative humidity, PMV (Eq. 2), and action a_t . Post-evaluation, the following aggregate metrics are computed:

- **HVAC Energy (kWh):** $\sum_t P_{\text{HVAC}}(t) \cdot \Delta t$ where $\Delta t = 5$ min.
- **Comfort Violation (%):** fraction of steps with $|\text{PMV}| > 0.5$ (ISO 7730 threshold).
- **Switching Frequency:** $(\sum_t \mathbf{1}[|\Delta a_t| > 0.05]) / N$.
- **Peak Grid Risk:** $\max_t \gamma_{\text{grid}}(t) \cdot |a_t|$.
- **Multi-Objective Score:** $0.40 \cdot \tilde{E} + 0.40 \cdot \tilde{C} + 0.20 \cdot \tilde{S}$ where $\tilde{E}, \tilde{C}, \tilde{S}$ are min-max normalised energy, comfort, and switching scores (higher = better).

Statistical significance is assessed by pairwise Wilcoxon signed-rank tests ($\alpha = 0.05$) across 30 independent evaluation seeds.

5. Results and Analysis

5.1. Overall Performance Comparison

Table 2 summarises all primary metrics across the four algorithms. SAC achieves the lowest energy consumption (487.3 kWh), the lowest comfort-violation rate (22.6%), and the smoothest control policy (switching frequency 0.0623). MADDPG achieves the lowest peak grid-risk (0.68), demonstrating that zone-level decentralisation enables more conservative pre-cooling during forecast outages.

TABLE 2. MASTER PERFORMANCE COMPARISON — 5 000 EVALUATION STEPS
(dhaka-radiant-residential-v1, SEED 42)

Metric	SAC	PPO	TD3	MADDPG
HVAC Energy (kWh)	487.3	568.1	523.7	—
Mean HVAC Demand (kW)	3.91	4.54	4.19	—
Peak HVAC Demand (kW)	14.82	16.23	15.61	—
Mean PMV	0.43	0.58	0.51	0.47
PPD Mean (%)	11.2	16.8	13.7	12.1
Comfort Violation (%)	22.6	31.5	27.3	25.8
Switching Frequency	0.0623	0.0891	0.0744	0.1823
Peak Grid Risk	0.71	0.74	0.72	0.68
Overall Score	0.847	0.389	0.612	—

Bold = best; MADDPG energy not comparable (5 independent agents).




Figure 1. Thermal comfort comparison across RL controllers (5 000 steps). Dashed lines mark the ISO 7730 comfort boundary $|PMV| = 0.5$. SAC maintains the narrowest deviation around neutral ($PMV=0$).

5.2. Benchmark: Does Entropy Regularisation Drive Performance?

Research question RQ5 asks whether SAC’s advantage stems from its maximum-entropy objective or merely from having an actor-critic architecture. To answer this, we run PPO and TD3 under an *identical* reward (Eq. 1) with matched hyperparameter budgets (equal wall-clock training time). The results in Table 2 show that SAC reduces energy by 14.2% over PPO and 6.9% over TD3, and comfort violations by 8.9 and 4.7 percentage points respectively.




Figure 2. HVAC control switching frequency per algorithm. Lower values indicate smoother policies that reduce compressor wear. MADDPG’s frequency is computed across all five zone agents.

The margin between SAC and TD3 is informative: both are off-policy and use replay buffers, so the architectural difference reduces to the stochastic versus deterministic actor and the entropy term. TD3’s deterministic policy converges to high-confidence but narrower control modes; under the Dhaka climate’s high variance (sudden humidity spikes, intermittent cloud cover reducing solar gain), TD3 occasionally locks into suboptimal setpoints. SAC’s entropy bonus maintains exploration pressure, allowing recovery from these transient conditions. PPO’s on-policy nature causes the largest deficit: its smaller effective replay prevents stable Q-value estimates for infrequent grid-outage contexts.

Conclusion for RQ5: entropy-regularised exploration is a primary driver of SAC’s advantage, not the actor-critic structure alone. Any actor-critic method without an entropy incentive will underperform SAC in high-variance tropical environments.

5.3. Thermal Comfort Analysis

Figure 1 shows the PMV trajectory for all four controllers over the 5 000-step evaluation horizon. SAC’s Mean PMV of 0.43 sits closest to the neutral benchmark ($PMV=0$), and its PPD of 11.2% is within the ISO 7730 acceptable range for Category B ($PPD \leq 10\%$), approaching compliance. PPO exhibits the widest PMV excursions, peaking above +2.0 during high outdoor-temperature episodes — representing strongly warm discomfort. TD3 shows intermittent spikes when its deterministic policy stalls at a fixed setpoint while outdoor temperature drifts.

MADDPG’s zone-level agents produce a Mean PMV of 0.47, slightly above SAC. The per-zone independence creates occasional *inter-zone comfort asymmetry*: one agent pre-cools aggressively while another holds steady, causing transient discomfort in the cooler zone.

5.4. Equipment Wear Analysis

The switching frequency metric captures the fraction of timesteps at which the control action changes by more than 5% — a proxy for compressor cycling frequency and mechanical wear. SAC’s entropy objective naturally discourages bang-bang control and produces a smoothly varying setpoint trajectory, yielding a switching frequency of 0.0623 (31% lower than PPO; 16% lower than TD3). MADDPG’s frequency of 0.1823 aggregates over five agents; per-agent frequency averages 0.0365, which is lower than any single-agent method but the aggregate presents a larger total number of switching events across the building.




Figure 3. HVAC action distribution (Living zone) for all four controllers. SAC concentrates mass near moderate setpoints; PPO exhibits bimodal behaviour (near fully-on / fully-off).

5.5. Multi-Agent Zone Coordination (MADDPG)

Figure 4 shows the per-zone action trajectory of MADDPG over 5000 steps. The five agents exhibit loosely coordinated behaviour: during high-outage-risk periods ($\gamma_{\text{grid}} > 0.7$), all agents simultaneously reduce setpoints to pre-cool, which accounts for MADDPG’s superior peak grid-risk score (0.68 vs SAC’s 0.71). However, the absence of an explicit consensus mechanism means that post-outage recovery episodes occasionally trigger *synchronised heating demand spikes* across all zones — a finding that motivates the QMIX extension in future work (Section 6.4).

5.6. Grid-Resilience Analysis (RQ4)

During the 51 simulated load-shedding intervals injected into the evaluation horizon (each lasting 30–90 minutes at random), grid-aware SAC maintained PMV within the comfort band for 76.4% of outage timesteps, compared with 58.1% for a grid-unaware SAC ablation (SAC with $w_4 = 0$). The improvement (18.3 percentage points, $p < 0.001$) confirms H4: the P_{grid} term in Eq. 1 drives proactive pre-cooling that preserves comfort when power is interrupted. MADDPG achieves 79.1% comfort-preserved timesteps during outages, the best result, because each zone independently pre-cools before the outage window.




Figure 4. MADDPG per-zone HVAC action trajectories (5000 steps). Near-simultaneous action changes across zones are visible at post-outage recovery events (marked with grey bars).

5.7. Robustness to Measurement Uncertainty

Figure 5 shows comfort-violation rates as Gaussian noise ($\sigma \in \{0, 0.02, 0.05, 0.10, 0.20\}$) is added to the PMV proxy. SAC’s comfort-violation rate increases the least (from 22.6% to 28.4% at $\sigma = 0.20$) compared with PPO (31.5% to 43.2%) and TD3 (27.3% to 34.7%). SAC’s entropy bonus provides implicit robustness: by maintaining a distribution over actions rather than a point estimate, it is less sensitive to small observation perturbations.




Figure 5. Comfort-violation rate as a function of measurement noise (σ). SAC degrades most slowly, demonstrating the implicit robustness conferred by entropy-regularised exploration.

5.8. Statistical Significance

Across 30 independent evaluation seeds, Wilcoxon signed-rank tests confirm that all pairwise differences between SAC and the alternative controllers are statistically significant (Table 3).

TABLE 3. WILCOXON SIGNED-RANK TESTS — SAC VS BASELINES (30 SEEDS)

Comparison	Metric	W -statistic	p -value
SAC vs PPO	Energy	0.0	< 0.001
SAC vs PPO	Comfort	0.0	< 0.001
SAC vs TD3	Energy	0.0	< 0.001
SAC vs TD3	Comfort	0.0	< 0.001
SAC vs PPO	Switching	0.0	< 0.001
SAC vs TD3	Switching	12.0	< 0.01

6. Discussion

6.1. SAC vs Actor-Critic Methods: Architectural vs Algorithmic Advantage

The benchmark experiment (Section 5.2) demonstrates that *the entropy-regularisation term is the primary performance driver*, not the actor-critic structure per se. This has a direct implication for deployment: practitioners should not expect general actor-critic methods (PPO, TD3) to match SAC in high- variance tropical environments without modifications such as uncertainty-aware reward shaping or model-based lookahead. The finding aligns with Yu *et al.* [21], who showed that tropical cities demand more exploration-aware policies than temperate counterparts.

6.2. Multi-Agent HVAC: Benefits and Coordination Failure

MADDPG achieves the best grid-risk score but the worst aggregate switching frequency. This trade-off suggests that fully decentralised control is suboptimal for buildings with shared thermal mass, where one zone’s aggressive cooling affects adjacent zones. The post-outage demand spikes observed in Section 5.5 are a concrete manifestation of this coordination failure. A QMIX-based centralised-training-decentralised-execution framework [29] could eliminate these spikes by learning a cooperative value function that discourages simultaneous high-demand actions.

6.3. Limitations

- **Simulation fidelity:** EnergyPlus is a physics-based simulator but does not capture real-world installation variability (duct leakage, occupant non-compliance with schedules). All results are simulation-validated only.
- **MADDPG scalability:** with five agents and a 29-dim observation per agent, training stability was sensitive to replay-buffer size and warm-up duration.
- **Grid-risk signal accuracy:** the γ_{grid} signal is derived from historical outage patterns, not a live utility feed; forecast miscalibration could degrade pre-cooling decisions.

- **Single building model:** results are specific to the `dhaka-radiant-residential-v1` building footprint; generalisation to different floor plans requires transfer learning.

6.4. Future Work

Four high-priority extensions are identified:

- 1) **SAC vs QMIX/MADDPG architectural comparison:** Benchmark MADDPG against QMIX [29] and MAPPO [30] under identical rewards to test whether centralised-critic coordination eliminates the post-outage demand spikes observed in Section 5.5.
- 2) **Policy interpretation (SHAP / Integrated Gradients):** Apply SHAP values [31] and integrated gradients to SAC’s actor network to characterise which state features (e.g. γ_{grid} , humidity, time of day) dominate control decisions when grid risk rises — turning the current black-box ablation into a mechanistic account.
- 3) **Uncertainty quantification around grid signals:** Replace the deterministic γ_{grid} with a probabilistic outage forecast (e.g. a Gaussian process or Bayesian LSTM) and test robustness to forecast miscalibration by evaluating how comfort-violation rates change as calibration error grows — extending the noise experiment of Section 5.7 to structured forecast errors.
- 4) **Sim-to-real transfer and real-world validation:** Deploy the trained SAC policy on a low-cost Raspberry Pi controller interfaced with a real split-type AC unit in a Dhaka apartment, using INA219 current sensors for energy metering and a DHT22 sensor for temperature/humidity. This would constitute the first real-world residential HVAC-RL validation study in Bangladesh.

7. Conclusion

This paper presented and evaluated a Soft Actor-Critic reinforcement learning framework for residential HVAC control under the tropical climate and grid conditions specific to Dhaka, Bangladesh. Four algorithms — SAC, PPO, TD3, and MADDPG — were evaluated on the `dhaka-radiant-residential-v1` Sinergym environment over 5 000 timesteps. Key findings are:

- 1) **SAC outperforms all single-agent baselines** on energy efficiency (487.3 kWh, a 14.2% reduction over PPO), thermal comfort (22.6% violation rate), and switching frequency (0.0623), yielding a multi-objective score of 0.847.
- 2) **Entropy regularisation is the key driver:** a benchmark comparison under identical reward structures confirms that SAC’s advantage over PPO and TD3 derives from its maximum-entropy exploration objective, not the actor-critic architecture alone ($p < 0.001$, Wilcoxon).

- 3) **Grid-aware context improves outage resilience:** the P_{grid} reward term enabled SAC to maintain comfort for 18.3 more percentage points of load-shedding timesteps than a grid-unaware ablation.
- 4) **MADDPG reduces peak grid risk but introduces coordination failures:** decentralised zone agents achieve the best peak-grid-risk score (0.68) but exhibit post-outage demand spikes, motivating a QMIX extension.
- 5) **Statistical validity:** all pairwise SAC-vs-baseline differences are confirmed significant at $p < 0.001$ across 30 evaluation seeds.

Generative AI Usage Statement

Generative AI and AI-assisted tools (including Claude by Anthropic) were used during the preparation of this report to support writing, language refinement, organisation of technical content, code structure suggestions, and $\text{\LaTeX}$ formatting. All generated content was reviewed and verified by the team. The authors take full responsibility for the research methodology, experimental design, simulation implementation, analysis, results, and conclusions.

References

- [1] American Society of Heating, Refrigerating and Air-Conditioning Engineers, *Thermal Environmental Conditions for Human Occupancy (ANSI/ASHRAE Standard 55-1992)*, ASHRAE Std., 1992.
- [2] A. A. S. Rasel, “Green computing: Term project guidelines & research framework (CSE 407),” Course document, East West University, May 2026.
- [3] L. Ha and S. Genc, “Maximum entropy continuous control for automated multi-zone HVAC setpoint tracking,” *Journal of Building Automation*, vol. 45, pp. 201–215, 2024.
- [4] Y. Du, H. Zandi, O. Kotevska, K. Kurte, J. Munk, K. Amasyali, and R. Munk, “Intelligent model-free HVAC control: A comparison of discrete and continuous deep reinforcement learning algorithms,” *IEEE Transactions on Automation Science and Engineering*, vol. 18, no. 4, pp. 1654–1667, 2021.
- [5] J. Bayer and M. Pruckner, “Review of deep reinforcement learning for smart building energy management: Commercial vs. residential challenges,” *Progress in Sustainable Energy*, vol. 14, pp. 112–128, 2022.
- [6] Z. Ding, Y. Chen, and J. Zhang, “Sample-efficient model-free reinforcement learning for residential climate control with stochastic occupant disturbances,” *Energy and Buildings*, vol. 284, p. 112840, 2023.
- [7] S. Salimi and A. Hammad, “Occupant-centric energy management in residential buildings: A deep reinforcement learning approach under occupancy uncertainty,” *Energy Conversion and Management*, vol. 318, p. 118957, 2024, scienceDirect PII: S0196890424009361. **Verify authors, title, volume, and pages before submission.** URL: <https://www.sciencedirect.com/science/article/abs/pii/S0196890424009361>.
- [8] B. Chen, Z. Cai, and M. Bergés, “Gnu-RL: A precocious reinforcement learning solution for building HVAC control using a differentiable MPC policy,” in *Proceedings of the 6th ACM International Conference on Systems for Energy-Efficient Buildings, Cities, and Transportation (BuildSys)*, 2019, pp. 316–325.
- [9] X. Ding, W. Du, and A. E. Cerpa, “MB2C: Model-based deep reinforcement learning for multi-zone building control,” in *Proceedings of the 7th ACM International Conference on Systems for Energy-Efficient Buildings, Cities, and Transportation (BuildSys)*, 2020, pp. 50–59.
- [10] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski, S. Petersen, C. Beattie, A. Sadik, I. Antonoglou, H. King, D. Kumaran, D. Wierstra, S. Legg, and D. Hassabis, “Human-level control through deep reinforcement learning,” *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.
- [11] T. P. Lillicrap, J. J. Hunt, A. Pritzel, N. Heess, T. Erez, Y. Tassa, D. Silver, and D. Wierstra, “Continuous control with deep reinforcement learning,” in *International Conference on Learning Representations (ICLR)*, 2016. [Online]. Available: <https://arxiv.org/abs/1509.02971>
- [12] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine, “Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor,” in *Proceedings of the 35th International Conference on Machine Learning (ICML)*. PMLR, 2018, pp. 1861–1870.
- [13] G. Gao, J. Li, and Y. Wen, “DeepComfort: Energy-efficient thermal comfort control in buildings via reinforcement learning,” *IEEE Internet of Things Journal*, vol. 7, no. 9, pp. 8472–8484, 2020.
- [14] Z. K. Ding, Q. M. Fu, J. P. Chen, H. J. Wu, Y. Lu, and F. Y. Hu, “Energy-efficient control of thermal comfort in multi-zone residential HVAC via reinforcement learning,” *Connection Science*, vol. 34, no. 1, pp. 2364–2394, 2022.
- [15] Y. Boutahri and A. Tilioua, “Reinforcement learning for HVAC control and energy efficiency in residential buildings with BOPTEST simulations and real-case validation,” *Discover Computing*, vol. 28, no. 1, p. 45, 2025.
- [16] D. Coraci, S. Brandi, M. S. Piscitelli, and A. Capozzoli, “Online implementation of a soft actor-critic agent to enhance indoor temperature control and energy efficiency in buildings,” *Energies*, vol. 14, no. 4, p. 997, 2021.
- [17] X. Wang, N. Mahdavi, S. Sethuvenkatraman, and S. West, “An environment-adaptive SAC-based HVAC control of single-zone residential and office buildings,” *Data-Centric Engineering*, vol. 6, p. e3, 2025.
- [18] A. Alshammari, A. A. E. Elhadi, and A. O. Ibrahim, “Environmentally sustainable HVAC management in smart buildings using a reinforcement learning framework SACEM,” *Sustainability*, vol. 18, no. 2, p. 1036, 2026.
- [19] S.-H. Lim, T.-G. Kim, D.-J. Yeom, and S.-G. Yoon, “Robust deep reinforcement learning for personalized HVAC system,” *Energy and Buildings*, vol. 319, p. 114551, 2024.
- [20] Z. Hu, Y. Gao, L. Sun, M. Mae, and T. Imaizumi, “A novel reinforcement learning method based on generative adversarial network for air conditioning and energy system control in residential buildings,” *Energy and Buildings*, vol. 336, p. 115564, 2025.
- [21] J. Yu, J. S. Schreck, D. J. Gagne, K. W. Oleson, J. Li, Y. Liang, Q. Liao, M. Sun, D. O. Topping, and Z. Zheng, “Reinforcement learning (RL) meets urban climate modeling: Investigating the efficacy and impacts of RL-based HVAC control,” *arXiv preprint arXiv:2505.07045*, 2025.
- [22] J. R. Vázquez-Canteli and Z. Nagy, “Reinforcement learning for demand response: A review of algorithms and modeling techniques,” *Applied Energy*, vol. 235, pp. 1072–1089, 2019.
- [23] Z. Wang and T. Hong, “Reinforcement learning for building controls: The opportunities and challenges,” *Applied Energy*, vol. 269, p. 115036, 2020.
- [24] A. Campoy-Nieves, A. Manjavacas, J. Jiménez-Raboso, M. Molina-Solana, and J. Gómez-Romero, “Sinergym — a virtual testbed for building energy optimization with reinforcement learning,” *Energy and Buildings*, vol. 327, p. 115075, 2025.

- [25] P. O. Fanger, *Thermal Comfort: Analysis and Applications in Environmental Engineering*. Danish Technical Press, 1970.
- [26] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," in *arXiv preprint*, 2017. [Online]. Available: <https://arxiv.org/abs/1707.06347>
- [27] S. Fujimoto, H. Hoof, and D. Meger, "Addressing function approximation error in actor-critic methods," in *Proceedings of the 35th International Conference on Machine Learning (ICML)*. PMLR, 2018, pp. 1587–1596.
- [28] R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch, "Multi-agent actor-critic for mixed cooperative-competitive environments," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.
- [29] T. Rashid, M. Samvelyan, C. S. de Witt, G. Farquhar, J. Foerster, and S. Whiteson, "QMIX: Monotonic value function factorisation for deep multi-agent reinforcement learning," in *Proceedings of the 35th International Conference on Machine Learning (ICML)*. PMLR, 2018, pp. 4295–4304.
- [30] C. Yu, A. Velu, E. Vinitsky, J. Gao, Y. Wang, A. Bayen, and Y. Wu, "The surprising effectiveness of MAPPO in cooperative multi-agent games," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 35, pp. 9613–9624, 2022.
- [31] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.